

# Merged Deep Learning Roadmap (CampusX + Krish Naik)

## Foundations

- Introduction to Deep Learning
  - Deep Learning vs Machine Learning
  - History & Applications of Deep Learning
  - Biological Neuron vs Perceptron vs Artificial Neuron
  - Perceptron Intuition & Geometric View
  - Training a Perceptron (Perceptron Trick)
  - Advantages & Disadvantages of Perceptron
  - Problem with Perceptron (Linear Separability)
- 

## Artificial Neural Networks (ANN / MLP)

- MLP Notation & Intuition
- Forward Propagation in Neural Networks
- Backpropagation Algorithm (What, How, Why)
- Chain Rule of Derivatives
- Vanishing & Exploding Gradient Problem
- Activation Functions:
  - Sigmoid (versions 1 & 2)
  - Tanh
  - ReLU
  - Leaky ReLU, Parametric ReLU
  - ELU, SELU
  - Softmax (for multi-class)
  - Which Activation Function to use When?
- Loss Functions vs Cost Functions
- Regression Loss Functions (MSE, MAE, etc.)
- Classification Loss Functions (Cross Entropy, Hinge Loss, etc.)
- Which Loss Function to use When?

- Gradient Descent Variants:
    - Batch GD, Stochastic GD, Mini-Batch GD
    - SGD with Momentum
    - Nesterov Accelerated Gradient (NAG)
    - AdaGrad
    - RMSProp
    - Adam Optimizer
    - Exploding Gradient Fixes
  - Weight Initialization Techniques:
    - Xavier, He Initialization
    - What not to do?
  - Regularization:
    - L1, L2, Weight Decay
    - Dropout Layers
    - Early Stopping
    - Batch Normalization / Layer Normalization
  - How to Improve Neural Network Performance
  - Hyperparameter Tuning (Keras Tuner, Optuna basics)
- 

## **ANN Projects (Basic Hands-On)**

- Customer Churn Prediction (Classification)
  - Handwritten Digit Recognition (MNIST)
  - Graduate Admission Prediction (Regression)
  - Finding Optimal Hidden Layers & Neurons
  - ANN Training End-to-End (with Optimizer + Loss)
  - ANN Deployment with Streamlit
- 

## **Convolutional Neural Networks (CNN – Computer Vision Basics)**

- CNN Intuition & Biological Analogy (Visual Cortex)
- Understanding Images (RGB, Channels, etc.)
- Convolution Operation

- Padding & Strides
  - Pooling (Max, Min, Average)
  - Flattening & Fully Connected Layers
  - CNN vs ANN
  - CNN Architectures: LeNet-5, AlexNet (intro)
  - Backpropagation in CNN (Convolution + Pooling layers)
  - CNN Practical Project: Cat vs Dog Classification
  - Data Augmentation in CNN
  - Pretrained Models (VGG, ResNet, Inception basics)
  - Transfer Learning vs Fine Tuning
  - Visualizing CNN Filters & Feature Maps
  - What does a CNN see?
- 

## **Recurrent Neural Networks (RNN – Sequential Data)**

- Why RNN? RNN vs ANN
  - RNN Architecture & Intuition
  - Forward Propagation Through Time
  - Backpropagation Through Time
  - Problems with RNN (Vanishing Gradients, Long Dependencies)
  - Simple RNN with Keras (IMDB Sentiment Dataset)
  - Prediction from RNN Model
- 

## **LSTMs & GRUs**

- Why LSTM?
- LSTM Architecture:
  - Forget Gate
  - Input Gate & Candidate Memory
  - Output Gate
- Training Process in LSTM
- Variants of LSTM (Peephole, Stacked, BiLSTM, etc.)
- GRU Architecture & Intuition

- GRU vs LSTM
  - Bidirectional RNN/LSTM/GRU
  - Practical Projects:
    - Next Word Prediction with LSTM
    - Sentiment Analysis with RNN/LSTM
    - Streamlit Integration with RNN/LSTM
- 

## **Sequence-to-Sequence Models & Attention**

- Encoder-Decoder Architecture (Seq2Seq)
  - Problems with Seq2Seq (long dependencies)
  - Attention Mechanism:
    - Bahdanau Attention
    - Luong Attention
    - Self Attention
    - Scaled Dot Product Attention
    - Multi-Head Attention
    - Positional Encoding
  - Layer Normalization (vs BatchNorm)
  - Complete Encoder Architecture
  - Complete Decoder Architecture
  - Cross Attention
  - Masked Multi-Head Attention
  - Final Linear + Softmax Layer
- 

## **Transformers (Modern Deep Learning)**

- Why Transformers?
  - Transformer Basic Architecture
  - Encoder-Decoder Attention Working
  - Transformer Inference Process
  - Applications of Transformers (BERT, GPT intro)
  - Epic History of LLMs (CampusX extra)
-

# Final Outcome

By completing all these topics, you will have:

- **Core DL foundations** (ANN, CNN, RNN, LSTM, GRU, Transformers).
  - **Hands-on projects** (Churn, MNIST, Cat vs Dog, IMDB Sentiment, Next Word Predictor, etc.).
  - **Advanced topics** (Attention, Transformers).
  - **Deployment basics** (Streamlit from Krish Naik).
- 

This merged roadmap is **all of Deep Learning** (foundation + CV basics + NLP basics).

After this, you'll be ready to move into **specialized tracks**:

- **Computer Vision advanced** (Object Detection, Segmentation, GANs).
- **NLP advanced** (Transformers, BERT, GPT, etc.).
- **Practical DL frameworks** (PyTorch/TensorFlow).

# Deep Learning Master Roadmap (with Best Video per Topic)

---

## Foundations

- Introduction to DL → **CampusX** (Day 1–2, deeper explanation).
  - Perceptron & MLP basics → **CampusX** (Day 5–7, with intuition & math).
  - Forward Propagation → **CampusX**.
  - Backpropagation → **CampusX** (detailed derivation).
  - Why Perceptron fails? Linear separability → **Krish Naik** (short, crisp).
- 

## Activation Functions & Losses

- Sigmoid, Tanh, ReLU family → **CampusX** (deep intuition).
- Softmax (multi-class) → **Krish Naik** (quick + coding).

- Loss Functions (MSE, Cross-Entropy, etc.) → **CampusX** (theory) + **Krish Naik** (hands-on).
- 

## Optimization & Training Tricks

- Gradient Descent variants (SGD, Momentum, RMSProp, Adam) → **CampusX** (full theory + math).
  - Weight Initialization (Xavier, He) → **CampusX**.
  - Regularization (Dropout, BatchNorm, Early Stopping) → **CampusX** (concepts) + **Krish Naik** (practical demo).
  - Hyperparameter Tuning → **Krish Naik** (he shows Keras Tuner).
- 

## ANN Projects

- Churn Prediction (Classification) → **Krish Naik** (Udemy hands-on).
  - MNIST Handwritten Digit Recognition → **CampusX** (step-by-step).
  - Graduate Admission (Regression) → **Krish Naik**.
  - ANN Deployment with Streamlit → **Krish Naik**.
- 

## CNN (Computer Vision Basics)

- CNN Intuition + Convolution/Pooling → **CampusX** (deep math explanation).
  - CNN Architectures (LeNet-5, AlexNet) → **CampusX**.
  - Backprop in CNN → **CampusX**.
  - Cat vs Dog Classification Project → **Krish Naik** (practical + easy to code).
  - Data Augmentation, Transfer Learning → **CampusX**.
  - Visualizing CNN filters & feature maps → **CampusX**.
- 

## RNN (Sequential Data)

- RNN Intuition + Architecture → **CampusX** (Day 40+).
  - Vanishing Gradients in RNN → **CampusX**.
  - Simple RNN Project (IMDB Sentiment) → **Krish Naik** (short demo).
- 

## LSTM & GRU

- LSTM Intuition + Gates (Forget/Input/Output) → **CampusX** (best for math + animation).

- Variants of LSTM → **CampusX**.
  - GRU Intuition & Comparison → **CampusX**.
  - Next Word Prediction with LSTM → **Krish Naik** (he codes this simply).
  - Sentiment Analysis with LSTM/GRU → **Krish Naik**.
- 

## Attention & Seq2Seq

- Encoder-Decoder Seq2Seq → **CampusX**.
  - Bahdanau Attention → **CampusX**.
  - Luong Attention → **CampusX**.
  - Self-Attention & Scaled Dot-Product Attention → **CampusX**.
  - Multi-Head Attention → **CampusX**.
  - Positional Encoding → **CampusX**.
- 

## Transformers

- Transformer Architecture (Encoder + Decoder) → **CampusX** (multiple lectures).
  - Masked Multi-Head Attention → **CampusX**.
  - Applications (BERT, GPT intro) → **CampusX** (Epic history of LLMs).
- 

# Final Formula

- **CampusX for concepts, math-heavy parts, attention/transfomers.**
- **Krish Naik for quick coding projects + deployment.**